Dr. James Price

University of Bristol /
GW4 Alliance

Isambard: tales from the world's first Arm-based production supercomputer

'Isambard' is a new UK Tier 2 HPC service from GW4

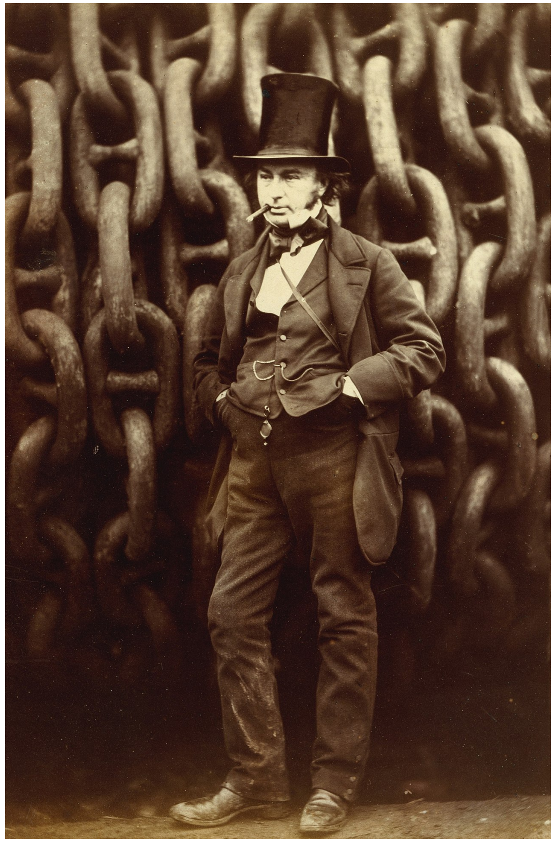

Isambard Kingdom Brunel
1804-1859

The tiered model of HPC provision

Tier 0: international

Tier 1: national

Tier 2: regional

Tier 3

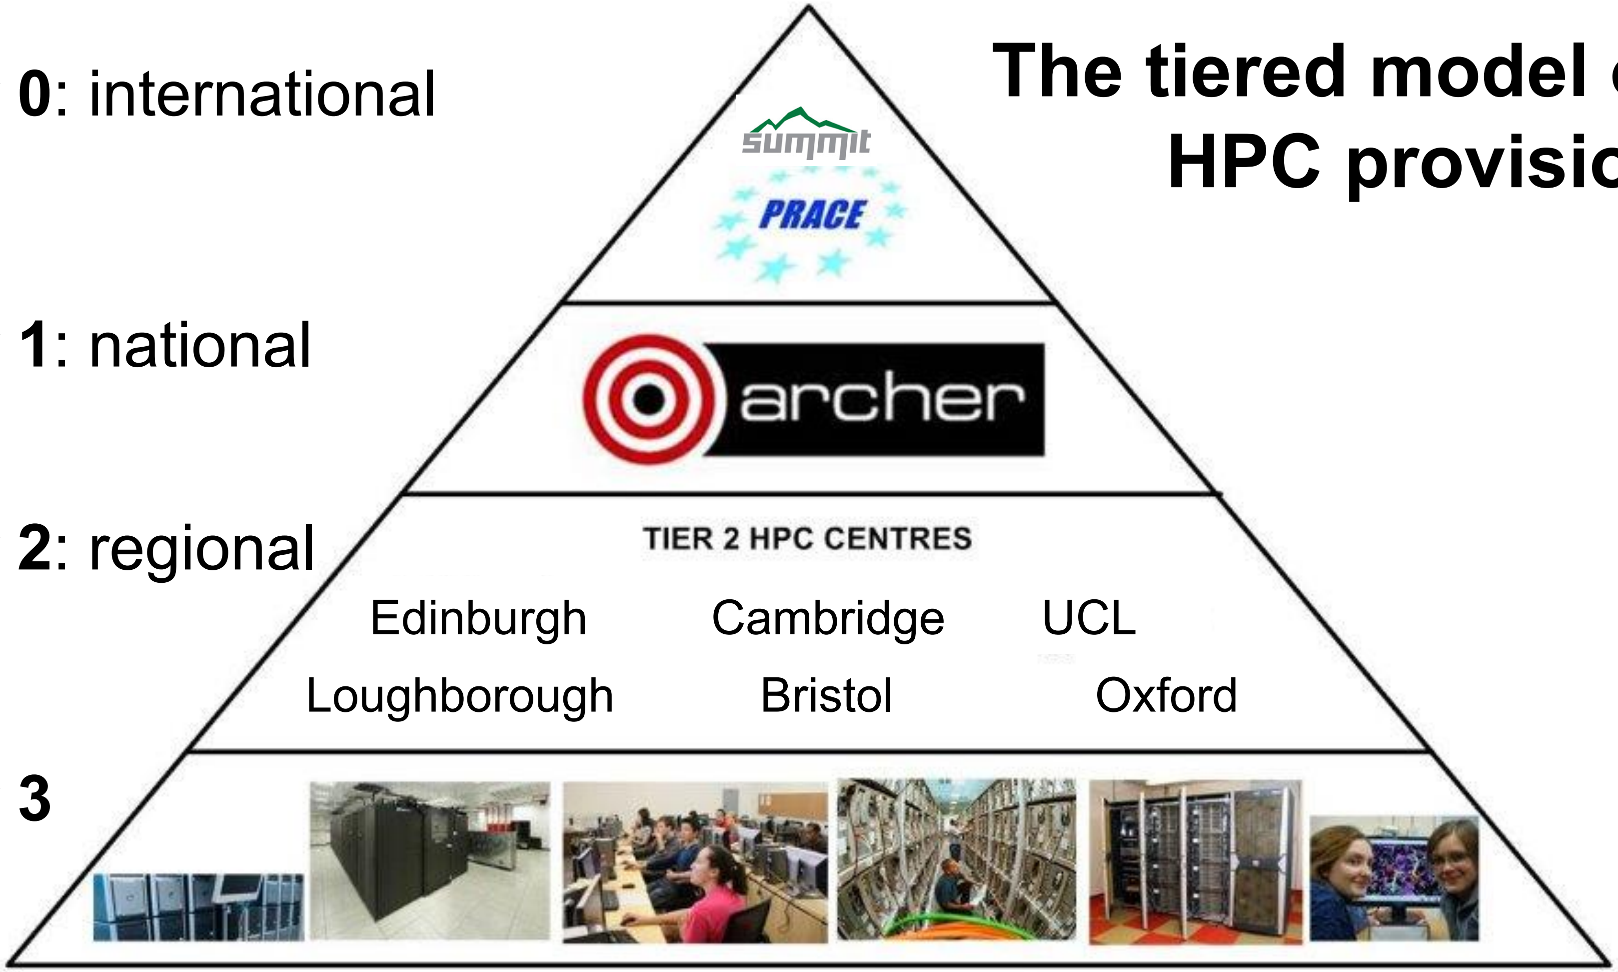

University of BRISTOL

EPSRC

GW4

Isambard system specification

- **10,752** Armv8 cores (168 x 2 x 32)
  - **Cavium ThunderX2 32 core 2.1GHz**
- Cray XC50 Scout form factor
- High-speed **Aries** interconnect
- Cray HPC optimised software stack
  - CCE, CrayPAT, Cray MPI, math libraries, …
- **Technology comparison:**
  - **x86, Xeon Phi, Pascal GPUs**
- Phase 1 installed March 2017
- Phase 2 (the Arm part) currently in bring-up
- £4.7m total project cost over 3 years

University of BRISTOL

EPSRC

Cavium ThunderX2, a seriously beefy CPU

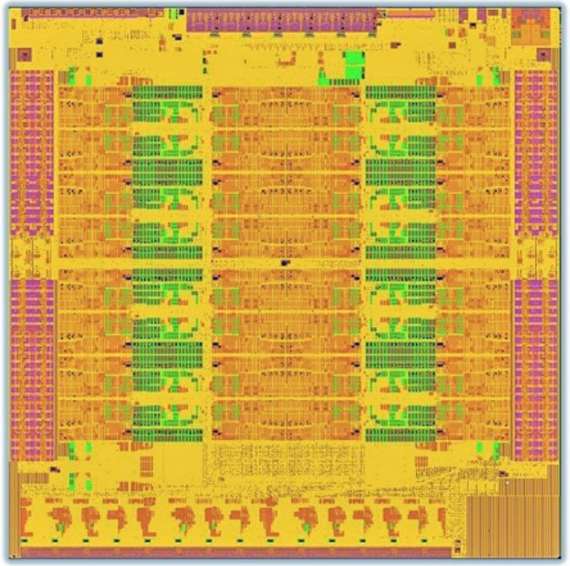

- 32 cores at up to 2.5GHz
- Each core is 4-way superscalar, Out-of-Order
- 32KB L1, 256KB L2 per core
- Shared 32MB L3
- Dual 128-bit wide NEON vectors
  - Compared to Skylake's 512-bit vectors, and Broadwell's 256-bit vectors
- 8 channels of 2666MHz DDR4
  - Compared to 6 channels on Skylake, 4 channels on Broadwell
  - AMD's EPYC also has 8 channels

University of BRISTOL

EPSRC

GW4

Benchmarking platforms

Processor	Cores	Clock speed GHz	TDP Watts	FP64 TFLOP/s	Bandwidth GB/s
Broadwell	2×22	2.2	145	1.55	154
Skylake Gold	2×20	2.4	150	3.07	256
Skylake Platinum	2×28	2.1	165	3.76	256
ThunderX2	2×32	2.2	175	1.13	320

BDW 22c Intel Broadwell E5-2699 v4, **$4,115** each (near top-bin)
SKL 20c Intel Skylake Gold 6148, **$3,078** each
SKL 28c Intel Skylake Platinum 8176, **$8,719** each (near top-bin)
TX2 32c Cavium ThunderX2, **$1,795 each** (near top-bin)

Key architectural comparisons (node-level, dual socket)

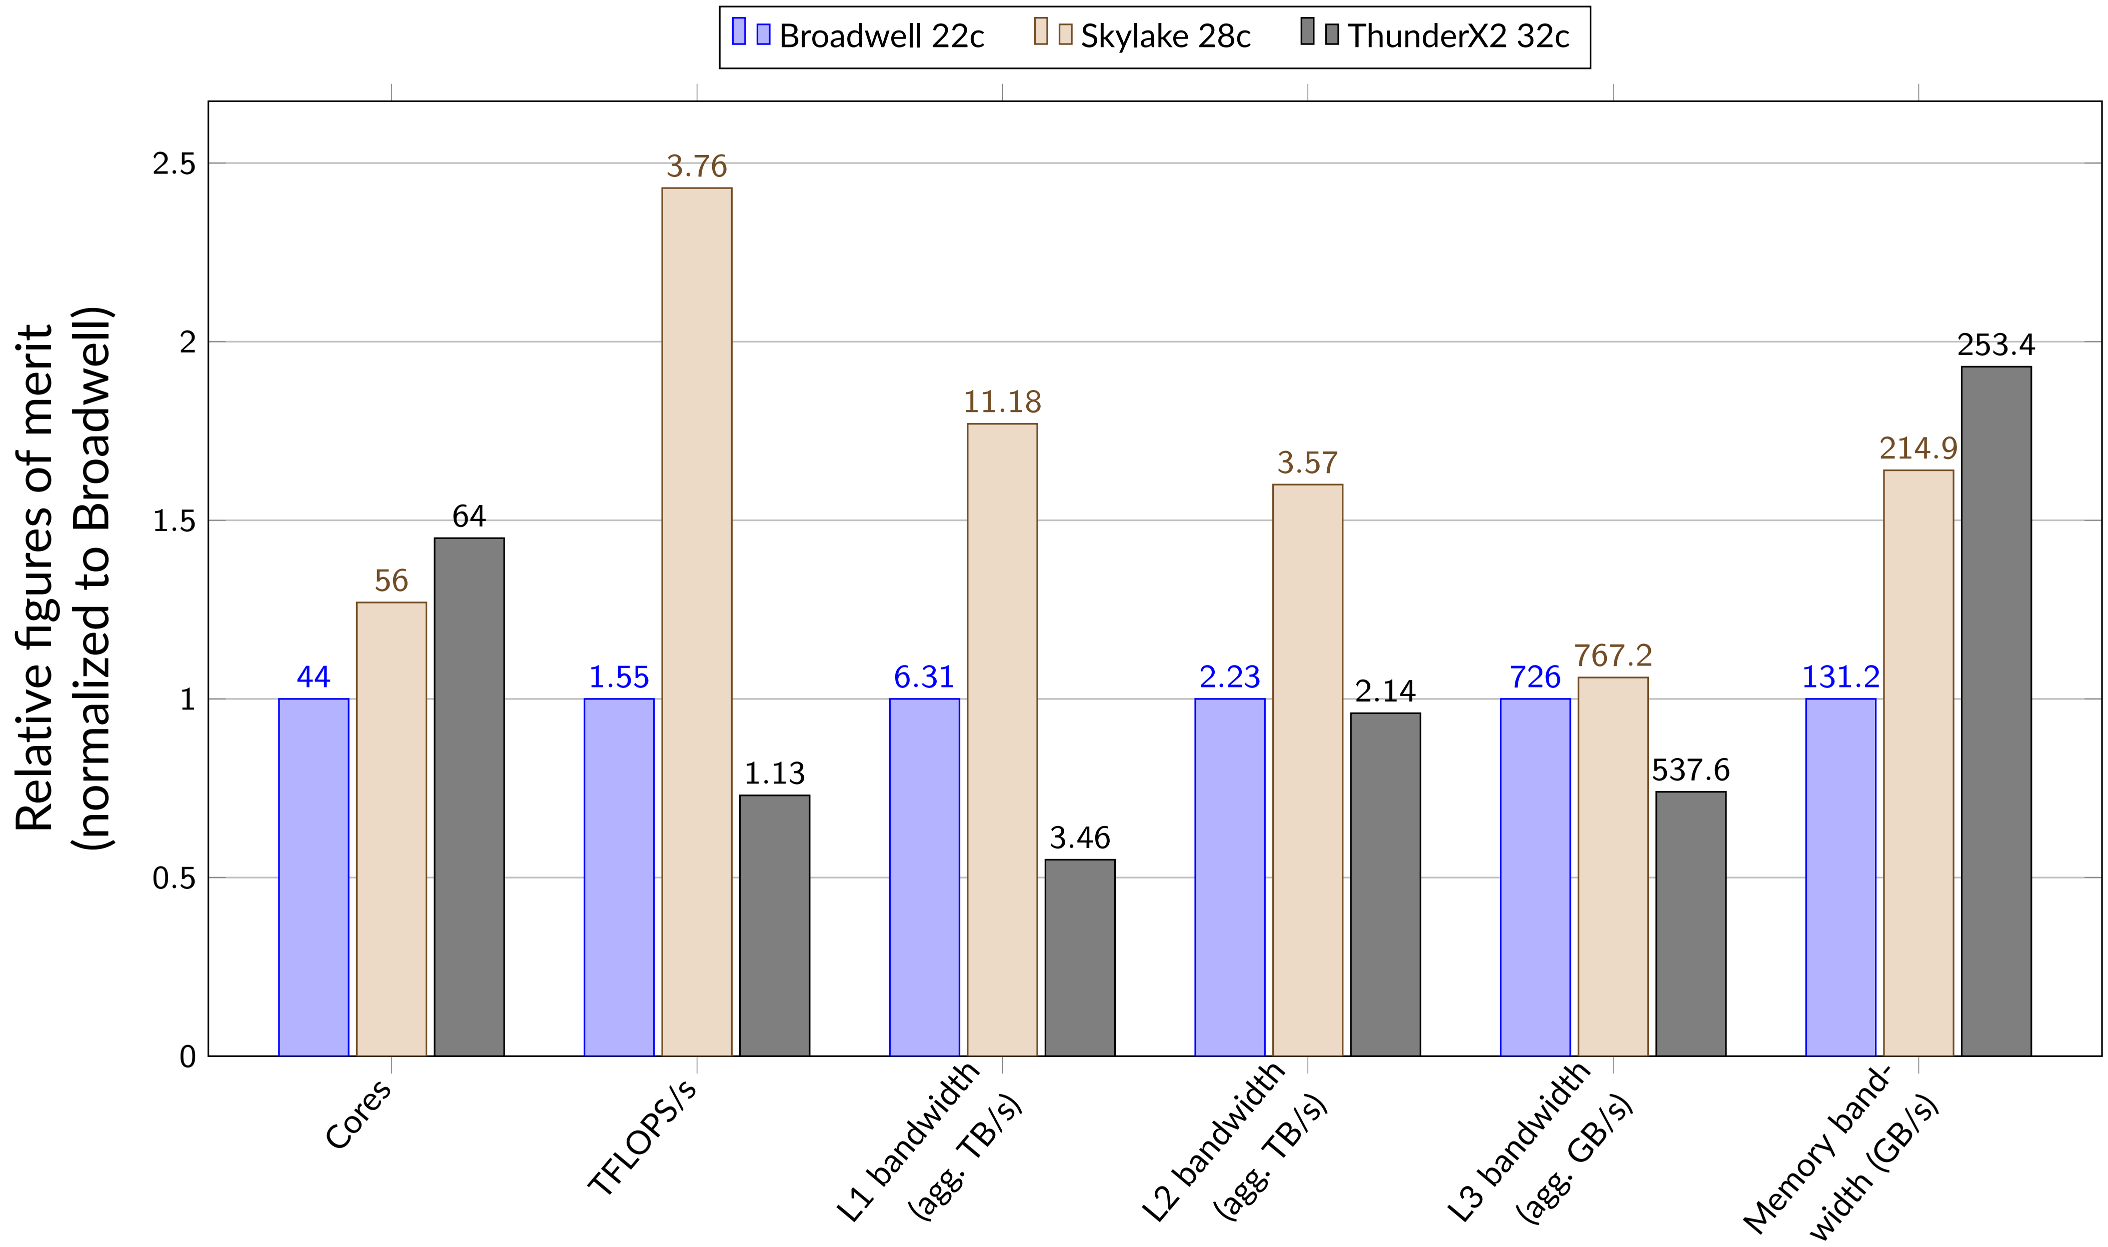

Isambard's core mission: deploying Arm in production HPC

Starting by porting/benchmarking/optimizing codes from the top 10 most heavily used on Archer:

- **VASP, CASTEP, GROMACS, CP2K, UM, NAMD, Oasis, SBLI, NEMO**
- Most of these codes are written in FORTRAN

Additional important codes for project partners:

- **OpenFOAM, OpenIFS, WRF, CASINO, LAMMPS, …**

Performance on heavily used applications from Archer

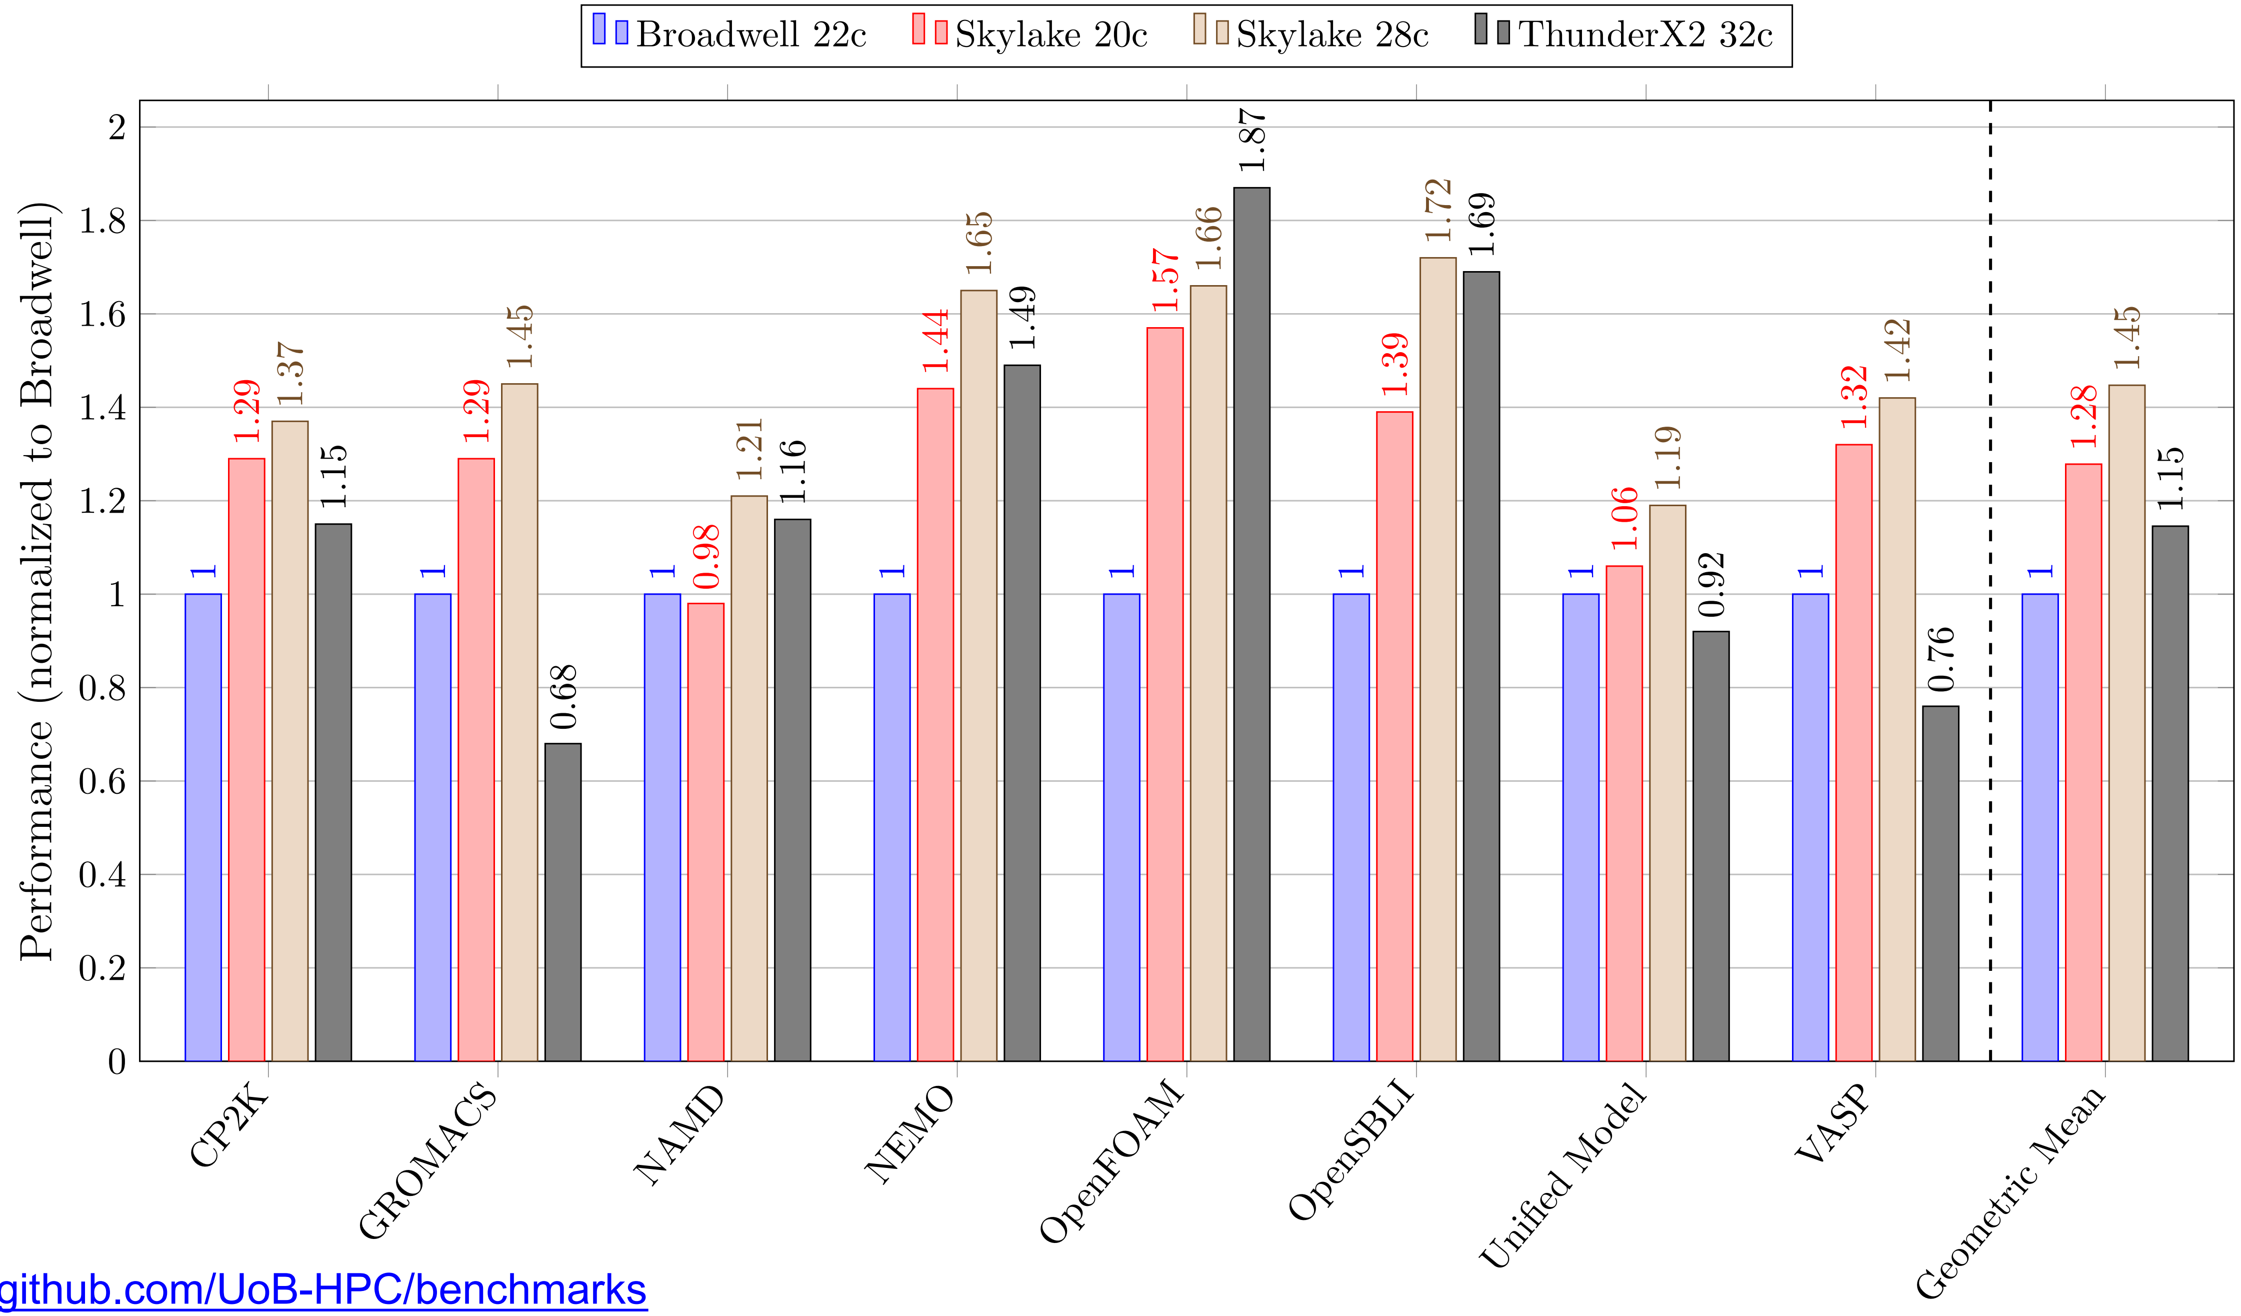

https://github.com/UoB-HPC/benchmarks

Performance summary

- Performance is competitive with contemporary Intel processors
  - ThunderX2 is **faster** when memory bandwidth is critical
  - ThunderX2 is **slower** when FLOP/s and L1 cache bandwidth matters
  - Even in the worst case, only drops ~30% performance versus Broadwell
- Next-gen Arm CPUs will increase FLOP/s + cache bandwidth
  - Introduction of SVE will allow vector width of up to 2048-bits
  - Fujitsu A64FX chip unveiled recently with 512-bit SVE
  - Expecting 512-bits to be a common choice for server chips

Benchmark	ThunderX2	Broadwell	Skylake
STREAM	Arm 18.3	Intel 18	CCE 8.7
CloverLeaf	CCE 8.7	Intel 18	Intel 18
TeaLeaf	CCE 8.7	GCC 7	Intel 18
SNAP	CCE 8.6	Intel 18	Intel 18
Neutral	GCC 8	Intel 18	GCC 7
CP2K	GCC 8	GCC 7	GCC 7
GROMACS	GCC 8	GCC 7	GCC 7
NAMD	Arm 18.2	GCC 7	GCC 7
NEMO	CCE 8.7	CCE 8.7	CCE 8.7
OpenFOAM	GCC 7	GCC 7	GCC 7
OpenSBLI	CCE 8.7	Intel 18	CCE 8.7
UM	CCE 8.6	CCE 8.5	CCE 8.7
VASP	GCC 7.2	Intel 18	Intel 18

https://github.com/UoB-HPC/benchmarks

Comparison of compilers on Arm

Exact same issues on x86

	GCC	Arm	CCE
STREAM	97%	100%	99%
CloverLeaf	92%	95%	100%
TeaLeaf	99%	95%	100%
SNAP	74%	87%	100%
Neutral	100%	94%	85%
CP2K	100%	BUILD	CRASH
GROMACS	100%	91%	CRASH
NAMD	83%	100%	BUILD
NEMO	-	-	100%
OpenFOAM	100%	97%	BUILD
OpenSBLI	-	-	100%
Unified Model	84%	72%	100%

https://github.com/UoB-HPC/benchmarks

Enabling co-design of future architectures with cycle-accurate simulation

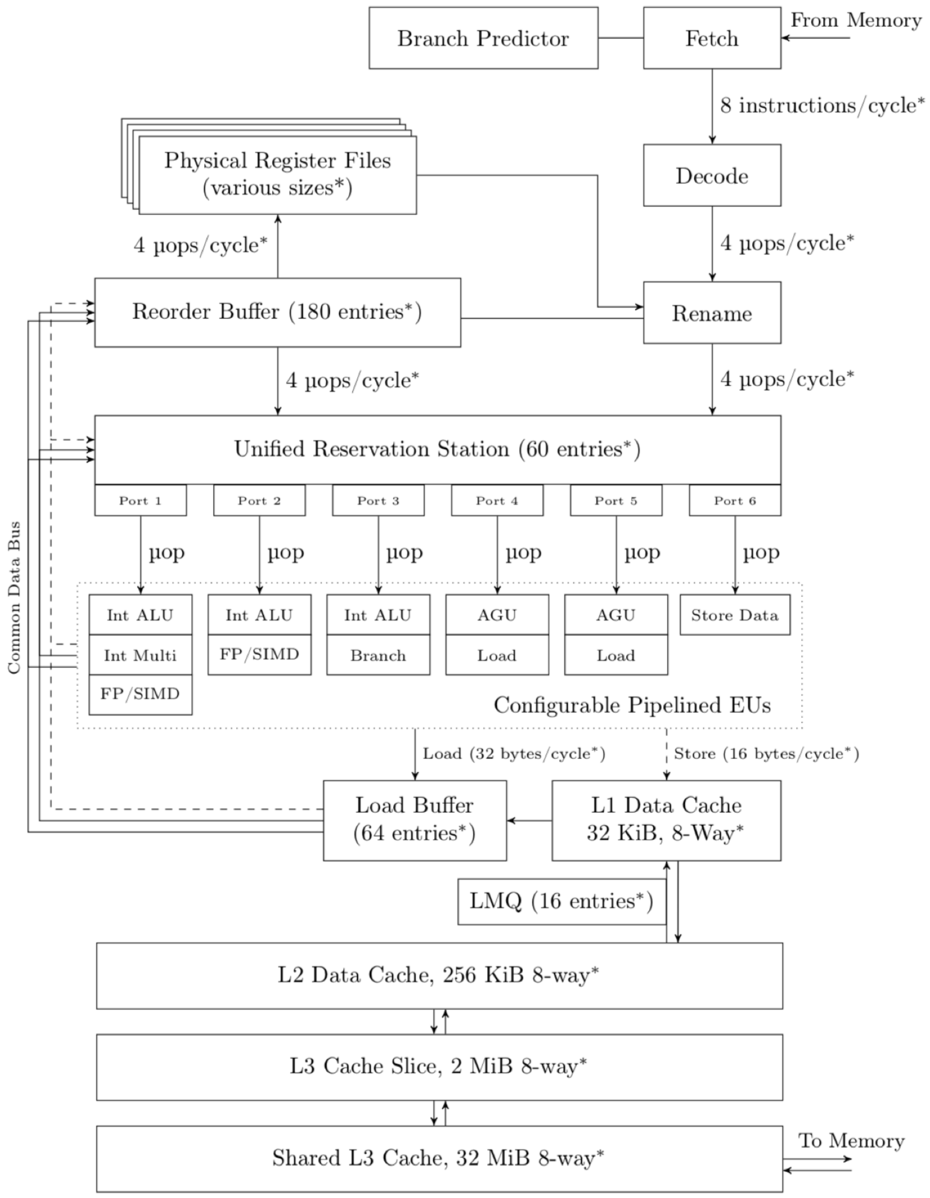

- We've developed a new configurable cycle accurate simulator in Bristol
- Within ~5-10% of TX2 hardware
- Highly configurable to almost any design of HPC CPU:
  - Planning a Post-K / A64fx version
  - Already supports SVE binaries
- Plan future support for x86, RISC-V, co-processors, …

EPSRC

GW4

Conclusions

- Results show **ThunderX2 performance is competitive with current high-end server CPUs,** while **performance per dollar is compelling**
- **The software tools ecosystem is already in good shape**
- The full Isambard XC50 Arm system is coming up now, we’re aiming to have early results to share at SC18
- The signs are that **Arm-based systems are now real alternatives for HPC,** reintroducing much needed competition to the market

For more information

Comparative Benchmarking of the First Generation of HPC-Optimised Arm Processors on Isambard

S. McIntosh-Smith, J. Price, T. Deakin and A. Poenaru, CUG 2018, Stockholm

http://uob-hpc.github.io/2018/05/23/CUG18.html

Bristol HPC group: https://uob-hpc.github.io/

Isambard: http://gw4.ac.uk/isambard/

Build and run scripts: https://github.com/UoB-HPC/benchmarks

Backup

Comparing performance per Dollar

- Hard to do this rigorously
  - RRP is not what anyone pays
  - Whole system cost has to be taken into account
  - Purchase price vs. TCO
- However, we *can* form some useful intuition
  - The following charts were generated by taking the performance results, dividing by the official published list prices of the CPUs only, then renormalizing to Broadwell

Performance per Dollar for applications